

EBU6505 – Reasoning and Agents

Block 1, Day 5 – Coding Agents and Knowledge-Grounded Reasoning

Dr Riasat Islam

Queen Mary University of London

Semester 1, 2026–2027
Teaching Week 3

What have we covered so far this week?

- MDP ideas and reward-based decision making for sequential problems
- Knowledge representation and graph structure
- Knowledge-graph reasoning for retrieval, memory, and explanation
- Knowledge graphs inside modern AI systems and agent workflows

Main idea: Day 5 closes Block 1 by asking one practical question: how do these ideas help when an AI system must act, not just answer?

What are we going to cover today?

- A simple synthesis of where knowledge graphs helped across this block
- What a coding agent is and how its workflow is organised
- Where graph structure helps code understanding, memory, and impact analysis
- What can go wrong, and how we reduce risk before Block 2

Main idea: Today is the bridge from knowledge-grounded reasoning to real agent systems that read, plan, act, and check.

- *Artificial Intelligence: Foundations of Computational Agents*, Poole and Mackworth (3rd ed.), Chapters 16–17
- *Artificial Intelligence: A Modern Approach*, Russell and Norvig (4th ed.), Chapter 10
- DeLong et al. (2025), *Neurosymbolic AI for Reasoning Over Knowledge Graphs: A Survey*
- Coding-agent platform links and safety references are collected at the end of this deck


Learning Outcomes for Today

By the end of this session, you should be able to:

- explain in simple terms what a coding agent does
- describe where a knowledge graph helps in a code task
- identify common failure modes in coding-agent systems
- connect Block 1 agent ideas to Block 2 uncertainty-aware decision models

Warm-up Question

Question: If an AI agent edits code for you, what must it know before it acts?



Main idea: A coding agent is useful only when it has the task, the right context, a way to check itself, and clear limits on what it may do.

Where Have Knowledge Graphs Helped This Week?

Question: What pattern keeps appearing across RAG, reasoning, vision, and generation?

RAG and search

graph paths keep retrieval close to the right entities and facts

Vision-language systems

graph links add context about objects and relations

Neuro-symbolic AI

graph structure adds explicit rules and traceable steps

Controlled generation

graph attributes help constrain what the system should create

Main idea: The graph keeps giving the same benefit: it turns isolated data into connected evidence that the system can follow.

When Is a Knowledge Graph Worth the Extra Work?

Question: Should we use a graph for every AI task?

Good fit

- many connected entities
- multi-hop questions
- need for explanation
- repeated reuse of the same facts

Weak fit

- one short document answer
- little relation structure
- no need for traceability
- graph maintenance costs more than it helps


Main idea: A graph helps most when relationships matter. If the task is simple and isolated, a graph may be unnecessary overhead.

What Is a Coding Agent?

Question: What makes a coding agent different from a chatbot that only answers questions?

It can do four things


- read code and docs
- plan a change
- edit files or run tools
- check whether the change worked



Main idea: A coding agent is an acting system. It reads context, changes things, and checks the effect.

What Context Must the Agent Collect First?


Question: What should the agent look up before it changes a repository?



Main idea: Good coding agents start by building context. Acting too early is how they break code, miss constraints, or violate policy.

Shared Architecture of Coding Agents


Question: What simple architecture appears in many coding-agent systems?



Main idea: The names differ across products, but the core pattern is similar: understand the task, gather context, act, then verify.

Planner–Coder–Tester–Reviewer Workflow


Question: Why do many systems split the work into roles instead of doing everything at once?



Main idea: Separating roles makes the workflow easier to inspect. Students should read this as a division of responsibility, not as magic intelligence.

One Reasoning Loop for a Code Change

Question: What does one full code-change loop look like in practice?



Main idea: Strong agents loop. They do not assume the first patch is correct; they revise after feedback from tests and checks.

Five Examples of Coding-Agent Products

Question: Which products should students recognise, even if the internal details differ?

ChatGPT Codex

repo edits
with checks

Replit Agent

build-run-fix loops
in one workspace

Z.ai GLM-5.1

long-horizon coding
in open-model workflows

Claude Code

long-context repo reading
and terminal iteration

Qwen Code

open workflows
with local control

Main idea: The shared pattern matters more than the brand name: gather context, use tools, revise, and check.

How Do the Products Differ in Practice?


Question: If the products follow a similar pattern, what actually changes from one system to another?

System	Typical strength	What to watch carefully
Codex	repo-grounded planning and patching	clear acceptance criteria still matter
Claude Code	strong large-repo reading and refactoring	command execution needs guardrails
Replit Agent	fast build-run-fix loop	secrets, dependencies, and deployment choices
Qwen Code	flexible local or open workflows	policy and permission setup is your job
Z.ai GLM-5.1	strong long-horizon coding	tool integration and workflow setup still matter

Main idea: Students do not need to memorize brand details. They need to understand the trade-off space: context depth, tool freedom, speed, and safety controls.

Where Can a Knowledge Graph Help a Coding Agent?


Question: What would the nodes and links mean in a software repository?



Main idea: A code graph is a knowledge graph over software. It connects files, functions, tests, APIs, owners, and policies in one structure.

Knowledge Graph for Impact Analysis


Question: If one function changes, what else might break?



Main idea: Graph links make impact analysis easier. The agent can see which APIs, tests, owners, and documents are connected to the change.

Knowledge Graph for Repository Memory


Question: How can the agent remember what happened earlier in a long task?



Main idea: Without memory, the agent forgets earlier findings. A graph lets it store task history in a reusable form instead of repeating the same work.

Failure Mode: Prompt Injection in a Repository

Question: What happens if a repository contains malicious instructions?



Safer response

- treat repo text as untrusted context
- keep user intent and policy above repo instructions
- gate dangerous actions behind checks or approval

Main idea: The agent should read the repository, not obey it blindly. Untrusted context is one of the easiest ways to trick a tool-using system.

Failure Mode: Secrets and Unsafe Commands

Question: Which two risks become serious as soon as the agent can use tools?

Secret leakage

- reads tokens from env files or logs
- pastes secrets into output or commits
- reduce risk with redaction and least privilege


Unsafe commands

- runs destructive or overbroad shell commands
- chains small steps into a risky action
- reduce risk with sandboxes and command gates

Main idea: Tool access is powerful, but that power creates risk. Permissions and safety checks are part of the agent design, not an optional extra.

Failure Mode: Passing Tests but Wrong Fix

Question: Why is 'tests passed' not always the same as 'the problem is solved'?



How to respond

use better coverage, stronger acceptance criteria, and human review for high-impact changes

Main idea: A shallow test suite can give false confidence. Verification is about the right checks, not only about having some checks.

Trust and Safety Checklist for Coding Agents

Question: What four controls should students look for in any serious coding-agent system?

1. Clear plan

task breakdown and acceptance criteria first

2. Permission tiers

low-risk reads are easier than risky writes or commands

3. Visible logs

keep a readable trace of what the agent saw and did


4. Human gate

pause when the task is ambiguous or high impact

Main idea: Do not ask only whether the agent can code. Ask whether it can code under controls you can inspect and trust.

Case Study: Enterprise Bug-Fix Incident


Question: How could a coding agent help on a real software incident without acting blindly?



Main idea: The useful role of the agent is to help map the problem, suggest changes, and support a safer validation flow.

Case Study: Multi-Agent Pull Request Review

Question: Why might one agent coordinate several specialist reviewers?



Main idea: Specialist agents can focus on one risk each. The coordinating agent combines their findings into a final decision.

What General Lesson Transfers Beyond Coding?

Question: Why does this discussion matter for healthcare, cyber, and finance too?

Healthcare

gather evidence
use policy
check risk first

Cyber

map systems
inspect alerts
test response safely


Finance

connect facts
justify decisions
keep an audit trail

Main idea: Coding agents are one example of a bigger pattern: useful agent systems need structure, evidence, and controls before they act in the real world.

Transition to Block 2 Applications

Question: What changes after Block 1?




Main idea: Block 1 focused on structured knowledge and grounded agent behaviour. Block 2 asks what to do when the world is uncertain and the system must still choose an action.

Key Takeaways

- Knowledge graphs help when the agent must follow relationships, not just read isolated text.
- Coding agents combine planning, retrieval, action, and verification in one loop.
- Graph structure can improve code understanding, memory, and impact analysis.
- Safety controls are essential because coding agents can read, write, and execute.

Main idea: The big lesson from Day 5 is simple: agent power comes from structure plus action, so trustworthy agents need structure plus controls.

In-Class Exercises



Use this time to work on today's in-class tasks and prepare one clear answer you can discuss.

Main idea: Use the exercise time to connect graph structure to a real coding-agent workflow and explain the safety value clearly.

Further Study

- Read Poole Chapters 16–17 and revisit AIMA Chapters 13–14 for agent structure and evaluation context.
- Practice: Poole Ch.16, Exercises 16.3 and 16.4.
- Practice: Poole Ch.17, Exercises 17.4, 17.5, 17.6, and 17.8.
- AIMA knowledge-representation exercise page: 7, 8, 16, 28, and 30.
- Focus on multi-step reasoning, structured evidence, and what must be checked before an agent acts.

Full References for This Lecture I

 Russell, S. and Norvig, P. (2021).
Artificial Intelligence: A Modern Approach (4th ed.).
Pearson. Chapters 13–14, pp. 459–540.

 Poole, D. L. and Mackworth, A. K. (2023).
Artificial Intelligence: Foundations of Computational Agents (3rd ed.).
Cambridge University Press. Chapters 9–10, pp. 381–454.

 DeLong, L. N., Mir, R. F. and Fleuriot, J. D. (2024).
Neurosymbolic AI for reasoning over knowledge graphs: A survey.
IEEE Transactions on Neural Networks and Learning Systems.
DOI: <https://doi.org/10.1109/TNNLS.2024.3420218>

Full References for This Lecture II



Mavromatis, C. and Karypis, G. (2024).

GNN-RAG: Graph neural retrieval for large language model reasoning.

arXiv preprint arXiv:2405.20139.

DOI: <https://doi.org/10.48550/arXiv.2405.20139>



Jiang, Y., Yang, Y., Xia, L. and Huang, C. (2024).

DiffKG: Knowledge graph diffusion model for recommendation.

In *Proceedings of the 17th ACM International Conference on Web Search and Data Mining*, pp. 313–321.

DOI: <https://doi.org/10.1145/3616855.3635850>



OWASP Foundation.

OWASP Top 10 for LLM Applications / GenAI.

<https://genai.owasp.org/llm-top-10/>

Full References for This Lecture III



OpenAI.

ChatGPT Codex product page.

<https://openai.com/codex/>



OpenAI.

Agents guide.

<https://platform.openai.com/docs/guides/agents>



Anthropic.

Claude Code overview.

<https://docs.anthropic.com/en/docs/claude-code/overview>



Replit.

Replit Agent documentation.

<https://docs.replit.com/replitai/agent>

Full References for This Lecture IV



Qwen Team.

Qwen Code documentation.

<https://qwenlm.github.io/qwen-code-docs/>



Z.AI.

GLM-5.1 overview and developer documentation.

<https://docs.z.ai/guides/llm/glm-5.1>